

GlobalGeoTree Dataset Analysis

Exploring source bias, temporal trends, and regional patterns in a global tree observation dataset

GlobalGeoTree team

2026-07-07

1 Introduction

Authors: Nataliia (source bias & Germany analysis), Anna (data structure & US analysis), Camille (worldwide distribution)

— Nataliia

This report analyses the GlobalGeoTree dataset, a global collection of approximately 6 million georeferenced tree observations covering species, taxonomic groups, and geographical locations of trees collected by diverse groups across the world. The goal is to understand how the dataset was assembled, what biases it contains, and how tree recording patterns vary across time, geography, and data sources.

GlobalGeoTree aggregates tree observations from a large number of independent sources: citizen-science platforms, institutional herbarium and museum collections, systematic field inventories, compiled floristic atlases, and various research and conservation projects. The loaded dataset covers observations from 2015 to 2024.

Because GlobalGeoTree is assembled from many independent sources, it does not reflect a systematic global sampling design — it reflects where people happened to record trees. The result is a dataset with strong geographic, taxonomic, and temporal biases that are source-driven rather than biologically meaningful. Understanding these biases is a prerequisite for any downstream analysis. To make the source structure interpretable, each observation’s free-text source name is classified into one of seven groups using keyword matching (see the Methodology section for details and limitations).

The report is organised in three parts. **Part I** examines global patterns: which sources dominate the dataset and how they shape what we see, how observations have changed over time, and how taxonomic composition and species richness vary across the world. **Part II** zooms into Germany as a whole and then Berlin in detail. Baden-Württemberg and NRW appear only through selected density and NUTS3 district maps that provide visual context for interpreting Berlin. **Part**

III examines the United States, the most-represented country in the dataset, analysing species composition and geographic patterns at the national level.

All charts in this report show **recording intensity** — how many observations were recorded and where — not a direct census of trees. High values can reflect actual tree presence, but they equally reflect where observers and platforms are most active.

Setup

Packages are loaded below; inline comments note their role in the report.

The following variables define the dataset year bounds and how many items appear in rankings and maps.

2 Methodology

This section summarises the analytical choices that the code implements. It provides context for interpreting the results without requiring the reader to examine every code block.

Data scope. The loaded GlobalGeoTree observation file covers the period 2015–2024. Rows with missing country codes or missing species names are removed during loading. The remaining records form the working dataset for the global, Germany, and US analyses.

Source classification. Each observation carries a free-text source field with thousands of distinct values. A keyword-matching function (`classify_source()`) assigns each source name to one of seven groups: iNaturalist, Pl@ntNet, Other citizen-science apps, GBIF (aggregator), Institutional & survey data, Reference & thematic databases, and Other / unknown. The rules are applied in order, and the first match wins. This approach works well for English-language source names and major platforms but can misclassify source names in other languages or those using project-specific abbreviations. The “Other / unknown” group is a catch-all for everything that did not match any rule.

Geographic boundaries. For the Germany analysis, administrative boundaries at the NUTS1 level (Bundesländer) and NUTS3 level (districts) are loaded from GISCO, the geographic information system of the European Commission. Berlin district boundaries come from a separate open-data project. If a local copy of these files exists, it is used directly; otherwise the file is downloaded from the public URL on first render.

Visualisation choices. Maps with individual observation points use reproducible random subsamples for readability and to keep the PDF size manageable (see below). Aggregate charts, counts, richness tables, rarefaction analyses, and source-analysis summaries always use the full relevant data. English common names for species are fetched from the GBIF API and cached locally.

Subsampling strategy. The full dataset contains over 6 million observations. Rendering every row as an individual map point is prohibitively slow and produces visually indistinguishable results compared to a well-chosen subsample. All point maps in this report therefore use reproducible random subsampling (fixed seed 42) via a `sample_up_to()` helper, which returns all rows when

the data is smaller than the requested cap. Worldwide non-source point maps draw up to 40,000 observations (per facet where applicable). Source-analysis maps draw up to 50,000 per source group. Density maps (kernel density estimation) also draw up to 50,000 observations, as KDE smoothing is stable well below this threshold. US maps follow the same approach with US-specific caps. Germany maps use smaller region-specific caps, though the German subset is small enough that most caps are not reached and all available observations are displayed. Crucially, all analytical summaries — counts, rankings, richness calculations, and rarefaction analyses — use the complete, unsampled data.

Reading density maps. Several maps in this report use smooth filled density contours (2D kernel density estimation) to highlight where observations are most concentrated. The legend shows contour levels from low to high: darker or more saturated colours indicate areas where more observations are clustered per unit area. These are *relative* density levels within each individual map — the same level number on two different maps does not imply the same absolute observation count. As with all maps in this report, density maps show *recording intensity*, not true tree abundance.

Part I: Global Patterns

3 Data Preparation

3.1 Data Loading and Cleaning

The raw CSV is read with `fread()` for speed. String fields are trimmed, numeric columns are type-cast, and rows with missing country codes or species names are removed. Each observation's free-text source name is classified into one of seven source groups (see Methodology). Continent labels are derived by spatially joining observation coordinates to a world countries shapefile, with manual corrections for Russia (split at 60°E) and Réunion (reassigned to Africa).

Measure	Value
Observation file	data/GlobalGeoTree.csv
Dataset period	2015–2024
Rows analysed	5,991,220
Distinct sources	2,462
Distinct source groups	6
Distinct countries	211
Distinct species	20,747
Germany rows	204,231

The table above confirms that the dataset loaded correctly and gives an overview of its scope. The analysis uses 5,991,220 observations in total.


3.2 Dataset Overview

— Anna

Column	Description
sample_id	Entry ID
country_code	Code of the country (e.g. US for the USA)
level0	Broad vegetation category grouping plants by shared ecological and physiological traits (e.g. Evergreen Broadleaf)
level1_family	Taxonomic Family — a group of related genera sharing common evolutionary characteristics (e.g. Asphodelaceae)
level2_genus	Taxonomic Genus — a group of closely related species (e.g. Aloe)
level3_species	Taxonomic Species — the most specific classification level (e.g. Aloe africana)
location	Country name
source	Source from which the occurrence record was obtained (e.g. iNaturalist Research-grade Observations)
species_key	Species identifier
year	Year of the observation
longitude	Geographic longitude
latitude	Geographic latitude


3.3 Distribution of the data across the world

Data distribution across the different countries:



We see that most of the data come from developed western countries, and especially the US that have a very important number of observations compared to the other countries.

Yearly distribution of the data in different countries: is there any evolution of the data collected throughout the years? Do some countries present large imbalance in the data distribution? How could this affect our analysis?



We see on this graph that the observations across the different countries were collected during a period of 10 years and that a slight increase in the number of yearly observations has been seen in the last years, from 2022 to 2024.

4 Source and Platform Bias

— *Nataliia*

4.1 Source Classification

Each observation's source field is matched against keyword rules to produce one of seven groups (ordered from most to least specific):

- **iNaturalist** — the largest citizen-science biodiversity platform worldwide.
- **Pl@ntNet** — a plant-identification app with a stronger presence in French-speaking countries and the Global South.

- **Other citizen-science apps** — regional portals (e.g., Artportalen, iRecord, Naturgucker, NatureMapr).
- **GBIF (aggregator)** — re-distributes records from thousands of organisations; may *overlap* with any other group.
- **Institutional & survey data** — herbarium/museum collections and systematic field campaigns.
- **Reference & thematic databases** — compiled floristic atlases, global plant databases (e.g., BIEN, RAINBIO, Tropicos), Wikidata, and invasive-species tracking.
- **Other / unknown** — everything that did not match any keyword rule.


The keyword rules work well for English-language names and major platforms but can misclassify multilingual or opaque source names; the “Other / unknown” group likely includes recognisable sources whose names did not match any rule.

4.1.1 Classification coverage

Of all 5,991,220 observations, **97.3%** are assigned to one of the six named groups, and **2.7%** fall into “Other / unknown”. The higher the coverage rate, the more reliably the source-group comparisons throughout this report represent the full dataset.

4.2 Source Concentration


iNaturalist alone contributes **64.3%** of all observations; the three citizen-science groups together account for **84.4%**.



The Pareto chart shows extreme source concentration: a handful of individual sources account for the vast majority of classified observations, and the cumulative share curve flattens quickly. This means that any systematic bias in the top few sources will propagate through the entire dataset.

Observation share by source group

Citizen-science platforms dominate. GBIF (aggregat




The bar chart confirms that citizen-science platforms collectively dominate the dataset, with iNaturalist as the single largest contributor. Institutional and reference sources, while individually smaller, cover regions and taxa that citizen-science platforms may miss.

4.3 What Is in “Other / Unknown”?

“Other / unknown” contains 1,458 individually named sources that did not match any keyword rule. The chart below shows the largest contributors within this heterogeneous catch-all group.

Largest individual sources inside

“Other / unknown” contains 1,458 differentl




Percentages are share of the Other / unknown group, not of the whole dataset.

4.4 Countries and Regional Coverage

Source composition differs by country

Top 15 countries by observations; bars sum to 100% within each country. GBIF (aggregator) is shown separately — its true source mix is unknown.




The stacked bar chart reveals that source composition varies sharply by country. Some countries are almost entirely supplied by a single platform (e.g. iNaturalist in the US), while others draw from a more balanced mix. Countries dominated by a single source inherit that platform's geographic and taxonomic biases.

Shannon entropy converted to “effective number of source groups” quantifies source diversity per country (1 = one source dominates).


How source-diverse are the largest country samples?

Effective number of source groups is low when one source dominates a country.



Countries with low effective source diversity (close to 1) are essentially single-source samples, meaning their apparent species composition and geographic coverage are shaped almost entirely by one platform's user base. Countries with higher diversity draw from multiple independent collection efforts, making their data more robust to any single platform's biases.

4.5 Sources and Taxonomic Groups




The source mix is not uniform across plant families. Some families are predominantly recorded through citizen-science apps, while others rely more heavily on institutional surveys or reference databases. Families that are primarily recorded by a single source type may have geographic or temporal gaps that mirror that source's coverage limitations.

4.6 Geographic Footprints of Data Sources

Geographic coverage of major source groups

Different source groups leave visibly different spatial footprints.



Points show recording locations, not species ranges. Points are sampled for readability.



The faceted map reveals strikingly different spatial footprints across source groups. iNaturalist provides the most geographically widespread coverage but is concentrated in North America, Europe, and Oceania. Pl@ntNet complements it with stronger presence in French-speaking Africa and parts of South America. Institutional and survey sources tend to fill in localised regions that citizen-science platforms miss. These differences mean that where the dataset is geographically strong or weak depends directly on which sources are active in a given region.

5 Temporal Development

— *Nataliia*

5.1 Observation Volume Over Time


Between 2015 and 2024, observations changed from **198,874** to **813,995**, an overall change of **309.3%**.




Year-to-year changes are calculated within the 2015–2024 dataset period.

The strongest growth occurs in **2018** (+228,796 observations), and the strongest decline in **2024** (-122,071).

5.2 Which Sources Drive the Temporal Trend?




The stacked area chart shows that the dataset’s overall growth is largely driven by the expansion of citizen-science platforms, particularly iNaturalist. When iNaturalist’s contribution grows, the total observation count rises accordingly, meaning the overall trend reflects platform adoption more than any change in tree populations.



Shares are calculated within each year of the 2015–2024 dataset period.


The year-by-year source share view shows whether the source mix is shifting over time. If one source group’s share is growing while others shrink, the dataset’s overall character is changing — even if total volume increases. A stable composition would suggest that all sources are growing at roughly the same rate.



The chart above highlights the single largest year-over-year jumps for each source group, making it easy to identify which platforms drove the biggest surges in data volume and when those surges occurred.

5.3 Geographic Coverage Over Time

This subsection examines whether the dataset's geographic reach has expanded over time by looking at the rate of new countries entering the dataset and by comparing observation locations at selected snapshot years.




The rate of new countries entering the dataset shows whether the geographic reach is still expanding. A declining rate suggests that most accessible countries are already covered and further growth is primarily a matter of densification rather than discovery of new regions.

The maps below show observation locations at four snapshot years, revealing whether records spread into new regions or simply become denser in existing ones.

Map snapshots of observations over time

Selected years show whether records spread or intensify in existing regions.




Comparing the four snapshot years shows that the geographic footprint of the dataset was already broadly established early in the period: North America, Europe, and parts of Oceania and South America are visible from the first panel. Later panels show a clear increase in point density within these same regions, but comparatively little expansion into previously unsampled areas such as central Africa, interior Asia, or the Amazon basin. This confirms that the dataset's growth over time has been primarily a matter of intensification rather than geographic expansion.

6 Taxonomic Composition and Geographic Patterns

6.1 Species observed in this dataset

— Anna

What are the different tree families referenced in this dataset?




Definition of each tree category:

- Evergreen Needleleaf: An evergreen needleleaf refers to woody vegetation that retains needle-like or scale-like foliage year-round
- Evergreen Broadleaf: An evergreen broadleaf is a plant or tree that has wide, flat leaves (rather than needles or scales) and retains its foliage year-round
- Deciduous Broadleaf: a diverse group of flowering plants (angiosperms) characterized by flat leaves that are shed annually before winter (Oak, Maple, Beech, Birch, and Poplar/Aspen)
- Deciduous Needleleaf : a unique group of conifers that, unlike their evergreen counterparts, shed all their needles annually in the autumn (rare trees)

We see here that a very important part of the observations in this dataset are deciduous broadleaf trees. It makes sense because most of the tree species spread across the world are part of this family. On the contrary, Deciduous needleleaf are rare and located in much more restricted areas.

What are the different taxonomic families represented in this dataset and in which proportion are they represented?





Again, here we see that the most represented taxonomic families in this dataset are also the most represented ones across the world, especially in North America, Europe and the north of Asia, where most of the data come from. However, as the source bias analysis above shows, this apparent representativeness should be interpreted with caution: the dataset reflects where observers and platforms are most active, not a systematic global survey.

6.2 Geographical distribution of the observations

— Anna


How are the trees distributed depending on the latitude?





We see on this graph that the observation density is much more important on the north hemisphere than in the south hemisphere, but this can be linked to the amount of land that differs between north and south hemisphere or the difference in the amount of observations between the two hemispheres.

Larger representations are also often used in ecological studies, because 5° latitudinal bands contain a satisfying amount of information without showing too much noise.




The same conclusions can be drawn as before.

We usually take 5° slices because it is small enough to find local tendencies but also big enough to avoid noise and too much non useful information and generally fine enough to reveal geographical patterns (climatic gradients, biodiversity, species distribution, etc.) without creating too much noise.

6.3 Geographical observations of species

— Anna


How does the latitude influences the different species observed? (Latitude and specific richness)



Species richness: simple count of the number of different species present in a specific biological community, landscape, or region. Much more important in the north hemisphere than in the south hemisphere. More important in temperate areas, much lower at the equator and near the poles.


We see here that the maximum species richness are very close in the two hemispheres, but that the distribution is broader in the northern hemisphere than in the southern hemisphere (it is also because trees grow at higher latitudes in the northern hemisphere than in the southern hemisphere). Similar environmental conditions can be found in the two hemispheres but they are not spread following the same pattern.

How does the longitude influences the different species observed (longitude and specific richness)?




We see here that areas with a lot of specific richness correspond to the continental areas, but not much more can be taken out of this graph.

What is the local specific richness across the world?




We see here that, in the areas presenting observations, the species richness does not present too important variations and is homogenous in most areas. By crossing this with the previous results, we can think that the species richness does not vary but that the species involved do vary across the world.

How are the different tree species distributed depending on the latitude? Here, we only check for the most represented species.




With this figure, we see that among the most represented species in this dataset, they present different world distributions. Most of them are mostly present in the northern hemisphere, but some are also present in the southern hemisphere (such as *Dombeya acrifolia*), and some are almost evenly present in the two hemispheres, such as *Cornus acuminata*.

Another representation could be a heatmap representing the species depending on the latitude:




Again, we see here that most species are present mostly in the northern hemisphere, but some can also be present in both hemisphere or only in the southern hemisphere or near the equator (*Guazuma umifolia*, a tree largely present in Latin America and really rare in other areas).


boxplot family x latitude



Same but on a map. Both maps below use the shared worldwide point-map subsample (see Methodology for subsampling details).



Spatial distribution of the 12 most abundant



On this visualization, we can see the distribution of the different taxonomical families across the world: - families are very widely spread across the world such as the Asteraceae: this family contains a very large amount of different species, that live in temperate areas. They can survive various non-extreme conditions. - families are observed in more restrictive areas, such as the families Ericaceae and Pinaceae: these families are very adapted to cold and unwelcoming places such as the tundra and boreal forests or places that are at a high altitude.

7 Worldwide Distribution of the Data

— Camille

How are observations distributed globally, and what does this reveal about sampling bias and taxonomic diversity? This section addresses four guiding questions:

1. Where are observations geographically concentrated, and does the coverage reflect ecological reality or observer activity?

2. How does the distribution of foliage types vary across latitudes and continents?
3. How does taxonomic richness compare across continents once differences in sampling effort are accounted for?
4. Has the spatial coverage of the dataset expanded or merely intensified over the years?

In this section, the aim is to understand how observations are distributed globally, both spatially and temporally. This distribution matters because it directly shapes how the rest of the dataset should be interpreted: an over-representation of certain regions may reflect sampling bias (the activity of contributors in the field) rather than the true state of global tree biodiversity. It is therefore essential to characterize where and when the data was collected, and to identify areas of high and low sampling density. We start by examining the overall geographical spread of observations across continents, before looking more closely at how taxonomic diversity (families, genus, species) is distributed in space. Finally, we analyze how sampling effort has evolved over time, both globally and by continent, to assess whether spatial biases have shifted across the years.

7.1 Pre-Processing

For this analysis, observations are grouped by continent rather than by country. With data spanning more than 210 countries, a country-level analysis would fragment the dataset into categories too small and numerous to reveal meaningful patterns, while a continent-level grouping provides a small number of interpretable regions, each large enough to support robust comparisons of taxonomic diversity and temporal trends. The continent for each observation was derived by spatially joining the point coordinates (longitude, latitude) with a world countries shapefile (rnatuarearth). This approach, however, relies on administrative continent boundaries, which do not always align with the geographical logic relevant to this study. Two adjustments were therefore necessary:

- Russia is administratively classified as a single entity, but spans two geographical continents. Observations were reassigned to Europe or Asia based on the conventional Ural Mountains boundary (60°E longitude).
- Réunion Island, despite being a French overseas territory, is geographically located in the Indian Ocean off the African coast, and was therefore reassigned to Africa rather than being grouped with other European territories.

Observations for which no continent could be determined, as well as those located over open ocean (e.g., isolated island territories not covered by the country shapefile), were excluded from the dataset.

The resulting dataset, with the added continent column, is previewed in Table 3.

To ensure visual consistency throughout this section, each continent is assigned a fixed color, shown in Figure 1. This color scheme is derived from the Okabe-Ito palette, chosen for its colorblind-friendly properties, and is applied consistently across all subsequent maps and charts to facilitate comparison between figures.

[!h]

sample_id	country_code	level0	level1_family	level2_genus	level3_species	source	year	continent
1	MW	Evergreen Broadleaf	Acanthaceae	Anisotes	Anisotes formosissimus	iNaturalist Research-grade Observations	2024	Africa
2	MZ	Evergreen Broadleaf	Acanthaceae	Anisotes	Anisotes formosissimus	iNaturalist Research-grade Observations	2018	Africa
3	MZ	Evergreen Broadleaf	Acanthaceae	Anisotes	Anisotes sessiliflorus	iNaturalist Research-grade Observations	2022	Africa
4	MZ	Evergreen Broadleaf	Acanthaceae	Anisotes	Anisotes sessiliflorus	iNaturalist Research-grade Observations	2015	Africa
6	KE	Evergreen Broadleaf	Acanthaceae	Avicennia	Avicennia spec	naturgucker	2023	Africa
7	NG	Evergreen Broadleaf	Acanthaceae	Avicennia	Avicennia spec	Biodiversity of Some Coastal Communities in Akwa-Ibom State, Nigeria	2021	Africa

Table 3: An overview of the first few rows of the dataset



Figure 1: Colour scheme used to represent the continents in the following figures

7.2 Global overview

This first subsection provides a global overview of where observations are located. The goal is twofold: first, to visualize the overall spatial coverage of the dataset and assess which regions of the world are represented; second, to quantify this coverage at the continent level in order to detect potential imbalances in sampling effort. Figure 2 projects a reproducible subsample of observations on a world map, colored by continent, offering a first visual impression of the dataset’s geographical spread without plotting millions of individual points. Figure 3 then complements this map with a quantitative summary, showing the total number of observations recorded per continent. Finally, Figure 4 goes one step further by highlighting areas of high and low observation density within continents, revealing local concentrations that a simple point map or continent-level count cannot capture on their own. For instance, distinguishing a continent with many observations spread evenly across its territory from one where the same volume of observations is concentrated in a few localized hotspots.




Figure 2: Geographical projection of observations by continent

Figure 2 shows a strongly uneven geographical coverage of observations. North America (green) and Europe (blue) display a dense, near-continuous coverage, closely following major population centers and road networks (a pattern typical of citizen science platforms like iNaturalist, where sampling effort is driven by observer density rather than tree distribution alone). Asia (orange) covers a considerable geographical extent, but coverage is markedly uneven: it is dense in East Asia, India, and along the Trans-Siberian corridor, yet visibly sparse across Central Asia and Siberia’s interior. Africa (yellow) shows a similar contrast, with observations concentrated in South Africa, East Africa, and coastal West Africa, while the Congo Basin and the Sahara remain largely unsampled. South America (light blue) is dominated by coastal Brazil and the Andean corridor, with the Amazon interior almost entirely absent from the dataset despite hosting one of the highest levels of tree biodiversity on Earth. Oceania (pink) is essentially limited to Australia’s populated coastline and a scattering of Pacific islands. Overall, this map already suggests that observation density mirrors human accessibility and citizen-science activity more closely than it reflects actual tree biodiversity (a hypothesis further examined in the following density map Figure 4).




Figure 3: Total number of observations by continent

Figure 3 quantifies the imbalance already suggested by Figure 2. North America and Europe together account for 74.4% of all observations (41.3% and 33.1% respectively), despite representing only a fraction of the world’s land area and tree biodiversity. Oceania follows at 10.3%, largely driven by Australia, before a sharp drop to South America (6.6%), Asia (5.5%) and Africa (3.3%).

This ranking reveals an important nuance compared to the previous map. Although Asia appeared visually extensive in Figure 2 (spanning a wide longitudinal range with a dense cluster of points), it accounts for only 5.5% of total observations, far below Europe or North America. This apparent contradiction is simply a matter of scale: individual points can appear numerous and spread out on a map while still representing a small share of a dataset containing several million observations overall. This illustrates why the two figures are complementary. The map shows where observations are located, while the bar chart shows how much each region actually contributes to the dataset.

The density map below is computed on a reproducible random subsample (see Methodology for the subsampling rationale).




Figure 4: Geographical density of observations

Note: the density estimate is based on a reproducible random subsample of 50,000 observations (drawn with the report-level map seed from the full dataset), rather than all 6.26 million records. Passing the full dataset to `geom_density_2d_filled()` caused a crash in ggplot2 3.5.2 due to memory limitations in the internal stat pipeline. At this sample size, kernel density estimation produces a visually identical result, as the smoothing operation is stable well below this threshold.

Figure 4 confirms and refines the pattern already observed in the previous two figures by revealing that observations are not only concentrated on certain continents, but highly clustered within a small number of localized hotspots. The most intense concentration is found in Western Europe (particularly around France, Germany and the United Kingdom), followed closely by the northeastern United States. Secondary hotspots appear along the U.S. West Coast, in Central America, in southeastern Australia, and around Taiwan / southern Japan, while the vast majority of the remaining land surface, including most of Africa, Asia and South America, shows minimal to no density. This concentration pattern reinforces the hypothesis raised earlier: sampling effort appears strongly correlated with population density and citizen-science activity rather than with the underlying distribution of tree biodiversity. Notably, even within high-observation continents such as North America (Figure 3), density is far from uniform. Most observations are concentrated along the East and West coasts, while large inland areas remain comparatively under-sampled.

Together, these three figures paint a consistent picture. The dataset’s geographical coverage reflects observer accessibility and activity far more than actual global tree distribution, a limitation that should be kept in mind, particularly when interpreting taxonomic diversity by continent in the following sections.

7.3 Spatial distribution of foliage types

Having established the overall geographical coverage of the dataset, this subsection shifts focus from where observations occur to what is being observed where. Each tree in the dataset is classified into one of four foliage types (Evergreen Broadleaf, Evergreen Needleleaf, Deciduous Broadleaf, or Deciduous Needleleaf). This classification closely tied to climate, since foliage strategy reflects how trees adapt to seasonal variation in temperature and water availability. The goal is therefore to examine whether these foliage types are distributed uniformly across the sampled regions, or whether they follow the climatic gradients expected from ecological theory, and to what extent this pattern can still be observed despite the sampling bias identified previously.

Three complementary perspectives are used to explore this question. Figure 5 first projects a reproducible subsample of observations for each foliage type, allowing their geographical range to be visualized directly on a world map. Figure 6 then aggregates this information into a quantitative comparison, showing how foliage composition varies from one continent to another. Finally, Figure 7 moves away from continental boundaries altogether, examining how the density of each foliage type is distributed along a continuous latitudinal gradient. This perspective is particularly relevant here, since climate (and therefore foliage strategy) is primarily structured by latitude rather than by administrative continental divisions.




Figure 5: Geographical distribution by type of foliage

Figure 5 reveals that Deciduous Broadleaf (orange) and Evergreen Broadleaf (green) share a broadly similar geographical footprint, both present across nearly all sampled continents and spanning a wide range of latitudes, from temperate to tropical regions. This similarity suggests that, at the continental scale visualized here, geographical presence alone does not clearly separate these

two categories. A finer, latitude-based view is needed to detect any underlying climatic gradient between them, which is examined in Figure 7. Evergreen Needleleaf (blue) is most heavily concentrated across the boreal and temperate zones of the Northern Hemisphere (North America, Europe, and northern Asia) but is also present, at lower density, in the southern parts of South America, southern Africa, and Oceania, suggesting an association with temperate conditions more broadly rather than strictly cold northern climates. Deciduous Needleleaf (pink) remains overwhelmingly concentrated at high latitudes, primarily across Siberia and parts of Canada, with only a lower-density secondary presence in the United States and Europe, confirming that, unlike Evergreen Needleleaf, this foliage type stays largely confined to high-latitude regions overall.




Figure 6: Geographical distribution by type of foliage

Figure 6 complements the previous maps by quantifying, for each continent, the relative composition of foliage types, independent of each continent's overall sampling volume. A clear contrast emerges between temperate and tropical/subtropical continents. North America and Europe share a similar profile, dominated by Deciduous Broadleaf (around 60–70% of observations), followed by Evergreen Broadleaf as the second-largest category (roughly 25–30%), with Evergreen Needleleaf making up a smaller but still notable share (around 10%), a share not found at comparable levels on any other continent, consistent with the boreal/temperate conifer forests characteristic of these regions. Oceania and Africa, by contrast, are clearly dominated by Evergreen Broadleaf, which accounts for the majority of observations in both continents (over 75% in Oceania, around 60% in Africa), with Deciduous Broadleaf as a smaller secondary category. South America and Asia present an intermediate profile, with Evergreen Broadleaf and Deciduous Broadleaf nearly evenly split at close to 50% each, suggesting that these two continents span a broader mix of climatic zones, from tropical to temperate, rather than being dominated by a single climate regime. Across all four of these continents, Evergreen Needleleaf drops to a marginal share compared to North America and Europe, and Deciduous Needleleaf is nearly absent everywhere except for a very thin presence in North America, Europe, and Asia, consistent with the highly restricted geographical range already observed in Figure 5. This continent-level breakdown reinforces the latitudinal pattern suggested by the maps: continents located predominantly at temperate latitudes (North America, Europe) show a foliage composition weighted toward deciduous and needleleaf strategies,

while continents dominated by tropical or subtropical land area (Oceania, Africa) are dominated by evergreen broadleaf species, with South America and Asia sitting in between due to their wider climatic range. This distinction by continent, however, remains an approximation, since large countries such as the United States or China span multiple climatic zones. This limitation is directly addressed by Figure 7, which examines this relationship along a continuous latitudinal axis rather than through continental boundaries.




Figure 7: Distribution of foliage types by latitude

Figure 7 confirms and refines, with much greater precision, the patterns suggested by the previous map and bar chart. Evergreen Needleleaf (blue) shows the most clearly defined latitudinal signal in the entire figure: a sharp, dominant peak around 45–50°N, with only a small secondary bump around 30–40°S, confirming that this foliage type is overwhelmingly associated with Northern Hemisphere temperate/boreal latitudes, with a much smaller presence in the Southern Hemisphere’s temperate zone. Deciduous Needleleaf (pink) is similarly concentrated at high northern latitudes, with a pronounced peak around 45–55°N and virtually no presence in the Southern Hemisphere. Its distribution is the narrowest and most latitude-restricted of the four foliage types, consistent with its very limited geographical range observed in Figure 5. Deciduous Broadleaf (orange) is dominated by a large peak around 45–50°N, overlapping closely with the two needleleaf types, alongside a much smaller secondary bump near 10°N and a minor bump around 30°S. Evergreen Broadleaf (green) stands out as the most latitudinally spread category, with density visible almost continuously from 40°S to 50°N, including a distinct concentration in the tropical band (0–20°N) that is more pronounced than for the other three types, plus a bump around 30–40°S and a peak around 45–50°N that overlaps with Deciduous Broadleaf. This last point directly resolves the ambiguity raised in Figure 5 and Figure 6: while Deciduous Broadleaf and Evergreen Broadleaf appeared to occupy similar geographical ranges at the continent level, this ridge plot

shows they are only partially overlapping in latitude. Evergreen Broadleaf is the only foliage type with a substantial tropical presence, while Deciduous Broadleaf, like both needleleaf types, remains overwhelmingly concentrated in the temperate Northern Hemisphere. This confirms that latitude (more than continental identity) is the primary driver structuring the geographical distribution of foliage types in this dataset, consistent with ecological expectations linking foliage strategy to climate.

7.4 Taxonomic richness by continent

This subsection turns from the spatial and climatic distribution of foliage types to a different question: “how much taxonomic diversity does each continent actually contain ?” This question, however, cannot be answered directly by simply counting the number of distinct species or genera observed per continent, precisely because of the sampling imbalance identified in the first subsection of this report. A continent with more observations is mechanically more likely to have recorded more distinct taxa, even if its underlying biodiversity is not necessarily higher. Comparing raw taxonomic counts between continents therefore risks conflating genuine biodiversity differences with differences in sampling effort alone.

Three complementary views are used to address this issue. Table 4 first reports the raw number of distinct taxa recorded per continent, providing a simple numerical summary of the dataset’s taxonomic content. Figure 8 then visualizes these same raw counts as a bar chart, making it easier to compare continents at a glance, while keeping in mind that, at this stage, these values remain directly influenced by each continent’s sampling volume, as shown in Figure 3. Finally, Figure 9 addresses this limitation directly through rarefaction curves, which estimate the expected number of unique species as a function of sample size, allowing continents to be compared on an equal footing regardless of how many observations were actually collected for each.

continent	Foliage	Families	Genus	Species	Complete_taxons
Africa	4	168	987	3831	3831
Asia	4	180	958	3809	3809
Europe	4	173	737	2437	2437
North America	4	203	1219	6528	6528
Oceania	4	175	886	4900	4900
South America	4	183	1099	4968	4968

Table 4: Table of the taxonomic richness by continent




Figure 8: Taxonomic richness by continent (raw values)

Table 4 and Figure 8 report the raw number of distinct taxa recorded per continent, at three taxonomic levels: family, genus, and species. All six continents show a comparable number of families (168–203), suggesting that the dataset captures a broadly similar range of major taxonomic groups everywhere. This similarity breaks down markedly at finer taxonomic resolution : genus counts range from 737 (Europe) to 1219 (North America), and species counts range even more widely, from 2437 (Europe) to 6528 (North America). North America stands out as the richest continent at every taxonomic level, while Europe consistently ranks lowest.

Comparing this ranking to the observation volume documented in Figure 3 reveals an interesting discrepancy. Europe, despite being the second most heavily sampled continent overall (33.1% of observations), records the lowest species richness of all six continents (even lower than Africa and Asia, which together account for less than 9% of total observations). Conversely, South America and Oceania, which together represent under 17% of observations, show species counts (4968 and 4900 respectively) roughly double that of Europe, despite Europe having been sampled nearly four times more intensively.

This pattern suggests that raw taxonomic richness, as shown here, is not simply proportional to sampling volume, but it does not yet establish whether these differences reflect genuine biodiversity patterns or are still partly shaped by uneven sampling effort within each continent. A question addressed directly by the rarefaction analysis in Figure 9.

Methodological note : Rarefaction curves are obtained through repeated random subsampling. For a given continent, a subset of n observations is drawn at random (without replacement) from all its recorded observations, and the number of distinct species present in that subset is counted. This process is repeated many times for each sample size n , and the results are averaged to produce

a smooth expected curve. This is repeated for increasing values of n , from a small sample up to the full number of observations available for that continent, producing one curve per continent.




Figure 9: Rarefaction curves by continent

Figure 9 delivers the central result of this subsection, and substantially revises the picture suggested by the raw counts in Figure 8. At the shared sampling effort marked by the dashed line ($n=198,000$, the size of Africa’s full sample), the ranking of continents changes considerably. Oceania and South America emerge as the most diverse, both reaching close to 4,000 expected unique species, closely followed by Africa (around 3,900). North America and Asia form an intermediate group (roughly 3,700 and 3,400 respectively), while Europe remains clearly the least diverse continent, with only around 1,300 expected species (less than a third of the leading continents, despite being the second most heavily sampled continent overall).

The most striking result of this analysis concerns North America. While it recorded by far the highest raw species count in Figure 8 (6,528 species), this advantage all but disappears once sampling effort is equalized. Its rarefaction curve is overtaken by Oceania, South America, and Africa well before reaching the comparison point. This confirms that North America’s apparent leadership in raw richness was largely an artifact of its considerably larger sample size (Figure 3), rather than evidence of genuinely higher tree diversity.

Europe presents the opposite and arguably more informative case. Its rarefaction curve remains consistently and substantially below all other continents across the entire range shown, including at very small sample sizes where all continents should, in principle, be more directly comparable. This indicates that Europe’s low species richness is not primarily a sampling artifact (as it was heavily sampled overall) but likely reflects a genuine ecological pattern, plausibly linked to lower tree species diversity in temperate European forests compared to the other regions.

Finally, the relative ranking of Africa, Oceania, and South America appears fairly stable and closely clustered across the observed range, suggesting a comparable level of underlying diversity between these three continents once sampling bias is accounted for.

7.5 Changes in the distribution over time

The two previous subsections examined the spatial and taxonomic distribution of observations as a static snapshot, without considering when this data was collected. However, given the strong reliance on citizen-science platforms such as iNaturalist, sampling effort is unlikely to have remained constant over time. It has likely grown alongside the adoption of these platforms, and may have expanded geographically as new contributors joined from previously under-sampled regions. This final subsection therefore introduces the temporal dimension, examining how both the volume and the geographical spread of observations have evolved over the years.

Figure 10 first tracks the total number of observations recorded each year, alongside the number of distinct countries contributing data, allowing global growth in sampling effort to be distinguished from genuine geographical expansion.

Figure 11 then breaks this trend down by continent, revealing whether this growth has been shared evenly across regions or driven primarily by a subset of continents, which would in turn affect the balance of the geographical biases identified earlier.

Finally, Figure 12 brings the temporal and spatial dimensions together, mapping sampled observations over successive two-year periods to directly visualize how the geographical coverage of the dataset has expanded, or remained concentrated, over time.

Taken together, these figures assess whether the spatial imbalances observed in the previous subsections are a stable, structural feature of the dataset, or whether they are progressively narrowing as data collection continues to grow.




Figure 10: Trends in data collection over time

Figure 10 reveals a marked and rapid growth in data collection, particularly from 2018 onward. After remaining relatively stable and low between 2015 and 2018 (150,000 – 200,000 observations per year), the number of observations increases sharply each year, peaking in 2022 at just over 1,000,000 annual observations, a roughly six-fold increase in only four years. This pattern is consistent with the growing global adoption of citizen-science platforms such as iNaturalist over this period, rather than a sudden change in tree biodiversity itself.

For his part, the number of distinct countries contributing observations (red line) shows an upward trend, growing from 166 in 2015 to over 198 by 2021, before stabilizing around 197–202 through 2024. This increase indicates a genuine geographical expansion of data collection over time, to the point where the dataset now includes at least one observation from nearly every country in the world (since there are only around 200 countries globally). This near-complete country coverage, however, should be interpreted carefully alongside the sharp rise in observation volume over the same period. The fact that almost every country is represented does not imply that sampling effort is evenly distributed within or between them. A single observation in a given country carries very different weight than the hundreds of thousands recorded annually in others, such as those in North America or Europe.

Whether this recent growth in volume has been accompanied by a more even distribution across continents, or remains concentrated in already-dominant regions, is examined more directly in the next figure (Figure 11).




Figure 11: Collection volume by continent and by year

Figure 11 confirms and refines the hypothesis raised in the previous figure. The overall growth in observation volume since 2018 is almost entirely driven by North America and Europe, while the other four continents remain comparatively flat throughout the entire period. North America shows the strongest and most sustained growth, rising from roughly 20,000 observations in 2015 to over 450,000 in 2023, before a slight decline in 2024. Europe follows a broadly similar trajectory until 2022, reaching nearly 385,000 observations, but then declines sharply over the following two years, dropping to around 220,000 by 2024. This pattern is not observed on any other continent. Asia, Africa, and Oceania, by comparison, remain at a much lower and comparatively stable level throughout the period, together never exceeding roughly 100,000–150,000 combined observations per year even at their 2022–2023 peak. This confirms that the massive increase in global volume documented previously has not translated into meaningful growth for these regions, and that the spatial imbalance identified earlier has, if anything, widened in absolute terms over the past decade. South America stands out with a distinct, isolated spike in 2018 (around 180,000 observations), sharply contrasting with its otherwise low and flat trend before and after that year. This one-off anomaly likely reflects a specific, temporary factor rather than a lasting shift in regional sampling activity, since volume returns to its previous low level as early as 2019.




Figure 12: Geographical scope of data collection by period (5 two-year periods)

Figure 12 provides a direct spatial confirmation of the trends identified in the two previous figures, while also introducing an additional dimension: foliage type (Evergreen Broadleaf, Evergreen Needleleaf, Deciduous Broadleaf). Across all five sampled period panels, the overall geographical footprint of the dataset remains remarkably stable. The same broad regions (North America, Europe, coastal South America, parts of Africa and Southeast Asia/Oceania) are already represented as early as 2015–2016, with no major new region appearing in later periods. This is consistent with the near-total country coverage observed in Figure 10: geographical breadth was established early, and subsequent growth in volume (Figure 11) has taken the form of densification within already-sampled areas rather than expansion into new ones. This densification is visually apparent when comparing early and recent periods: point coverage becomes progressively denser and more continuous over North America and Europe, particularly for deciduous broadleaf (orange) and evergreen broadleaf (green) foliage, while regions such as central Africa, the Amazon interior, and Central Asia remain visibly sparse throughout all five periods, showing little to no improvement over the decade.

Overall, this final figure reinforces the central finding of this subsection: the sampling bias identified earlier in the report is not a transient artifact of an early, immature dataset that has since been

corrected, but a structural and persistent feature that has remained stable over the full ten-year period covered by the data.

Part II: Germany and Berlin Analysis

— *Nataliia*

Part II examines tree observations within Germany. Germany was selected because it has relatively high observation density and contributions from multiple data sources, making it a good case for studying how global biases play out at the national and sub-national level. Berlin is examined in detail because this is where our team is based. Baden-Württemberg and NRW are included only as short map-based comparisons.

Table 5: Germany data check

item	value
Total observations (all countries)	5,991,220
Germany observations used in Part II	204,231

7.1 Species observed in Germany and common name lookup

To make charts easier to read, English common names are shown next to Latin names wherever possible. As described in the Methodology section, the GlobalGeoTree dataset contains `species_key`, a unique GBIF identifier for each species. This key is used to query the GBIF vernacular-names API once and cache the result in `data/external/species_common_names.csv`. On every subsequent render, the report reads from that file.

Table 6: German species overview

metric	value
Unique species in German observations	464
Species with a valid GBIF species key	464

Table 7: GBIF common names lookup result

metric	value
Species queried from GBIF	464
Species with an English common name	421

Species without an English name in GBIF are shown with their Latin name only in charts. The cached file can be deleted and re-generated at any time by deleting `data/external/species_common_names.csv` and re-rendering the report.

8 Germany: National Patterns

This section examines how tree observations are distributed across Germany as a whole: spatial concentration, temporal trends, source structure, and Bundesland-level variation.

8.1 Load Bundesländer boundaries

Bundesländer polygons are needed to draw state boundaries on Germany maps and to assign observations to federal states. The code first looks for a local GeoJSON file; if it is not found, the public GISCO NUTS1 file is downloaded.

Table 8: Bundesländer boundary check


item	value
Boundary status	Loaded Bundesländer / NUTS1 boundaries from data/external/germany_nuts1_2021.geojson
Number of Bundesländer polygons	16

8.2 Distribution of observations in Germany

Where are observations spatially concentrated across Germany? The smooth filled density map below provides a first overview of recording hotspots across the country.

Germany observations: smooth filled density

A cleaner hotspot view that is easy to read in a report.



This shows recording density in the dataset, not true tree abundance.


Darker areas show where more observations were recorded. These hotspots should be read as observation intensity rather than tree abundance: a location can appear dense because observers are active there, because a platform has strong coverage, or simply because access is easier.

8.3 Observations by year

How has recording activity changed over time in Germany? A dataset dominated by a few recent years may show spatial patterns that reflect current platform growth rather than long-term tree distribution.

Germany observations by year

Annual counts show when the Germany subset was most strongly represented




Peaks in this chart reflect peaks in recording activity, not necessarily more trees. A strong upward trend in recent years is common in citizen-science datasets as more users join the platforms.

8.4 Observations by source on the country map

The year distribution above shows *when* Germany observations were recorded. The next question is *who* recorded them — which platforms contribute the most records, and do they cover the country differently? Each panel shows where one source group collected its records.

Spatial footprint of major observation sources in Germany

Each panel shows sampled observation points from one source group.




If source panels look visually different, this is evidence of source bias: the national density map then reflects a mix of ecological signal and platform coverage.

8.5 Observations by source (bar chart)

The map above shows *where* each source records; this bar chart shows *how many* records each source contributes.

Germany observations by source group

The same grouping is used in the country map above.



The largest bar identifies the dominant source. A single dominant source can strongly shape the density patterns seen on all subsequent maps.

i Note

Two Pl@ntNet source names. Pl@ntNet observations and Pl@ntNet auto ID are both Pl@ntNet-derived but represent different validation routes: one is the regular shared-observation dataset, and the other contains automatically identified occurrences. They are therefore kept separate throughout this report.

8.6 Observations mapped to Bundesländer


Having seen the sources, we now look at where their records land across Germany's sixteen federal states. Observations are assigned to Bundesländer using a spatial join — each observation's coordinates are matched to the boundary polygon it falls within.

Table 9: Bundesland spatial join check

metric	value
Germany observations before join	204,231
Matched to a Bundesland	203,727
Bundesländer with observations	16


Germany observations by Bundesland

Observation counts are based on spatially joining coordinates



Observation frequency by Bundesland

Darker Bundesländer contain more Germany observations in



aries: GISCO NUTS1 regions. Counts: GlobalGeoTree Germany subset.

Together, the bar chart provides exact rankings while the map shows whether high-record states cluster in one part of Germany. Differences between Bundesländer reflect population density, platform coverage, and observer activity rather than differences in actual tree abundance. Nordrhein-Westfalen and Baden-Württemberg — the first and third largest Bundesländer by observation count — are revisited later as spatial comparison points for Berlin.

9 Berlin

Berlin is the city of our university, which is why it was selected as the main target for regional analysis in Germany. As a city-state, Berlin's entire area is urban, which may produce recording patterns quite different from larger, more rural federal states.

9.1 Load Berlin district boundaries

Berlin is handled separately because a finer boundary layer than Bundesländer is needed for city-level analysis. The district boundaries come from an open-data project and provide polygons for Berlin's local areas (Ortsteile).

Table 10: Berlin boundary check

item	value
Berlin boundary status	Loaded Berlin district/local-area boundaries from data/external/berlin_districts.geojson
Number of Berlin polygons	97

9.2 Berlin observations

Table 11: Berlin observation summary

metric	value
Berlin filtering method	Spatial join to Berlin district/local-area polygons
Berlin observations	5,192
Berlin recorded species	158
Berlin recorded genera	85
Berlin recorded families	43
Available source values	13
Available years	2015-2024


9.3 Berlin observation sources

Source composition is important context for every Berlin map. If one platform dominates, Berlin hotspots may partly reflect that platform’s recording behaviour rather than tree distribution.

Table 12: Berlin observations by source

source_short	observations	share	label
iNaturalist research-grade	3,133	60.3%	3,133 (60.3%)
Pl@ntNet auto ID	1,241	23.9%	1,241 (23.9%)
ArtenFinder	388	7.5%	388 (7.5%)
Pl@ntNet observations	151	2.9%	151 (2.9%)
naturgucker	108	2.1%	108 (2.1%)
Observation.org	101	1.9%	101 (1.9%)
EOD – eBird Observation Dataset	55	1.1%	55 (1.1%)
BoBO - Botanic Garden and Botani...	7	0.1%	7 (0.1%)
Xeno-canto - Bird sounds from ar...	4	0.1%	4 (0.1%)
7a3679ef-5582-4aaa-81f0-8c2545ca...	1	0.0%	1 (0.0%)
Database for alien invasive plan...	1	0.0%	1 (0.0%)
Database of field studies on env...	1	0.0%	1 (0.0%)
Spatially explicit database of t...	1	0.0%	1 (0.0%)

Berlin observations by source




9.4 Spatial distribution of Berlin observations

Which parts of Berlin show the highest recording density? Because Berlin is geographically compact, a single density map with overlaid district boundaries is sufficient to show spatial patterns.

Tree observation density in Berlin

Filled contours show observation density; local-area borders provide spa



The density map highlights which parts of Berlin have the most recorded tree observations. High-density areas are likely to correspond to parks, green corridors, and neighbourhoods with active observers rather than necessarily having more trees than other areas.

9.5 Berlin local areas: observations and species


Table 13: Top Berlin local areas by observation count

district	observations	species	sources
Grunewald	1,714	87	7
Wannsee	490	76	8
Dahlem	298	39	7
Zehlendorf	274	43	6
Hakenfelde	273	46	8
Marienfelde	261	49	6
Buch	247	60	7
Westend	165	32	5
Köpenick	163	45	7
Rahnsdorf	135	36	5


district	observations	species	sources
Nikolassee	127	43	7
Müggelheim	120	39	8
Blankenfelde	116	41	7
Tegel	114	29	8
Französisch Buchholz	109	35	5
Schlachtensee	92	22	4
Kladow	60	23	5
Gatow	45	24	5
Schmöckwitz	44	24	7
Heiligensee	40	15	5

Berlin observations by local area

Only the top 20 areas are shown.




Recorded species by Berlin local area



Observation frequency by Berlin local area

Square-root scaling keeps both high and medium-count areas visible.




The bar charts rank local areas by observation count and species count separately, while the choropleth map shows the spatial pattern. Local areas with both high observation counts and high species richness may genuinely be more biodiverse or may simply have more active recording communities.

9.6 Top recorded species in Berlin

Table 14: Top recorded species in Berlin

scientific_name	common_name_en	observations	share
<i>Cupania sylvatica</i>	Litchi family	750	14.4%
<i>Fagus sylvatica</i>	Beech	373	7.2%
<i>Quercus oblongata</i>	Acorn	258	5.0%
<i>Dryocopus martius</i>	Black Woodpecker	234	4.5%
<i>Prunus padus</i>	Bird Cherry	230	4.4%
<i>Euonymus europaeus</i>	Common Spindle	218	4.2%
<i>Schefflera macrophylla</i>	Aralia and Ivy family	180	3.5%
<i>Prunus serotina</i>	Black Cherry	169	3.3%
<i>Sorbus aucuparia</i>	European Mountain-ash	164	3.2%

scientific_name	common_name_en	observations	share
<i>Robinia pseudoacacia</i>	Black Locust	158	3.0%
<i>Carpinus betulus</i>	Caucasian Hornbeam	108	2.1%
<i>Croton laccifer</i>	Euphorbia family	107	2.1%
<i>Sambucus nigra</i>	Black Elder	97	1.9%
<i>Symphoricarpos albus</i>	Common Snowberry	97	1.9%
<i>Alnus glutinosa</i>	Alder	93	1.8%



The most frequently recorded species in Berlin reflect the city's urban tree stock, which includes both native species and commonly planted street and park trees. High observation counts for a species can indicate that it is genuinely common, that it is easy to identify and therefore often recorded by citizen-science users, or both.

9.7 Berlin oak tree groups

This section demonstrates a practical application of the dataset. Oak trees are a public-health concern in Berlin because the oak processionary moth (*Thaumetopoea processionea*) can cause allergic reactions and skin irritation in people and pets near affected trees. Mapping recorded oak locations can help identify areas where monitoring or public warnings may be needed.

Table 15: Berlin oak observation summary


metric	value
Berlin oak observations	258
Berlin recorded oak species	1
Share of Berlin observations that are oak	5.0%

Table 16: Top Berlin local areas by recorded oak observations

district	oak_observations
Grunewald	92
Dahlem	27
Hakenfelde	24
Buch	17
Wannsee	17
Rahnsdorf	15
Zehlendorf	15
Westend	9
Nikolassee	8
Tegel	5
Französisch Buchholz	4
Heiligensee	4
Köpenick	4
Müggelheim	4
Grünau	3
Kladow	2
Marienfelde	2
Schlachtensee	2
Schmöckwitz	2
Blankenfelde	1


Recorded oak tree groups in Berlin

Oak observations are filtered from Quercus records and mapped with l



Recorded oak observations by Berlin local area

This helps locate oak record groups that may need monitoring



The oak map and local-area ranking show where recorded oak groups are concentrated across Berlin. High-density areas in the dataset are candidates for targeted monitoring and public warning campaigns, though the records reflect observer activity and are not a complete oak census. This analysis demonstrates how the GlobalGeoTree dataset can be used for practical applications beyond general biodiversity research.

10 Baden-Württemberg and Nordrhein-Westfalen



The Berlin analysis above revealed that observations are heavily concentrated in a single location, raising questions about whether this pattern reflects a Berlin-specific data quality issue or is typical for German regions. To investigate, we compare Berlin with Nordrhein-Westfalen (NRW) and Baden-Württemberg — the first and third largest Bundesländer by observation count. For each state, a density map and a NUTS3 district choropleth show where observations are concentrated. Both should be read as recording intensity, not true tree abundance.

10.1 Baden-Württemberg

Baden-Württemberg is kept as a compact visual reference for Berlin. The two maps below show where observations are densest across the state and how observation frequency varies across NUTS3 districts.

Baden-Württemberg observations: smooth filled

A clean hotspot view for comparing broader concentration areas




This shows recording density in the dataset, not true tree abundance.

The filled density map is included only as a broad visual comparison with Berlin's observation-density pattern. It should be read as recording intensity, not true tree abundance.

10.1.1 Baden-Württemberg observations by NUTS3 district

Observation frequency across Baden-Württemberg

Square-root scaling improves readability when a few districts have





10.2 Nordrhein-Westfalen (NRW)

NRW serves the same role as Baden-Württemberg: a state-level comparison point for Berlin. The density map and NUTS3 district choropleth below follow the same layout.

NRW observations: smooth filled density

A clean hotspot view of where observations are concentrated inside




This shows recording density in the dataset, not true tree abundance.

As with Baden-Württemberg, this density map should be read as recording intensity rather than true tree density.

10.2.1 NRW observations by NUTS3 district

Observation frequency across NRW NUTS3 district

Square-root scaling improves readability when a few districts have v



Comparing the density maps of Baden-Württemberg and NRW with that of Berlin highlights a notable difference in spatial distribution. While observations in BW and NRW are spread much more evenly across their respective territories — with multiple distinct hotspots corresponding to major cities and populated areas — Berlin’s observations are heavily concentrated in a single location. This contrast reinforces the data quality concern raised earlier: the Berlin subset appears to reflect localised recording activity rather than a representative spatial sample of the city’s tree cover.

Part III: Analysis of US Data

— Anna

What species and geographic patterns characterise the most-represented country in the dataset? This section addresses two guiding questions:

1. Which species and tree categories are most frequently observed across the US?
2. How does species richness vary geographically across the US territory?


As with the worldwide analysis, US maps use reproducible subsampling for point and density visualizations while all aggregate analyses (richness, species counts) use the full US subset. See Methodology for details.

10.1 Filtration of the US data

10.2 Summary of the US data


	observations	species	genera	families
1	1875587	2982	662	155

10.3 Analysis of the species distribution in the US



The species showed as most observed ones are very common species in North-America (large-leaved beech for the first one). As before, we can say that the observed species in this dataset are, to some extent, representative of the reality in the US.

We now do the same but for the tree categories:




The distribution of the different tree categories is similar to the worldwide one, at least in terms of proportions. This distribution once again corresponds to the worldwide representation of the species.

10.4 Geographic distribution of the data

The density map below is computed on a reproducible random subsample (see Methodology for the subsampling rationale).

GlobalGeoTree observations in the US

Density estimate based on a subsample of 50,000 observations.




We can see on this map that the observations cover almost the whole US territory, but that areas are missing data, namely in the center of the US and in Alaska. This is coherent since these areas are less populated than the coastal areas. However, there are also those presenting the most trees in the US, especially in the national parks zones. The conclusions drawn from the other observations would then have to be put back into context.

Species Richness: - Depending on the latitude:

Species richness depending on the latitude

Number of distinct species recorded per 1-degree latitudinal band.




Here, we see that more species are observed in the northern part of the US (it is also the largest part of the US, while the south is more geographically restrained). This pattern may reflect both genuine ecological gradients and the concentration of sampling effort in certain regions. These hypotheses would need to be compared with real humidity and temperature data to disentangle biological patterns from recording bias.

Species richness depending on the longitude:


Species richness by longitude

Number of distinct species recorded per 1-degree longitudinal band.



Once again, we see that more species have been observed in the eastern part of the US, where more observations have been made. However, even if less observations have been made in the central part of the US these observations still show various species, which results in a lower species richness but still present (between -90 and -120°).


Local species richness across the US territory:



As we said before, the recorded species richness is higher near the coasts, where more observations have been made. The lower values in the interior may partly reflect lower sampling effort rather than genuinely lower diversity.

Geographic distribution of the 20 most represented species:

Spatial distribution of the 20 most observed species in the US




Displayed points: 50,000 sampled across the 20 most represented species.

We see that the most represented species can be spread: - all across the US territory: *Acer negundo* is the most common species of maple, which is a widely and famously spread tree in North America, - in specific parts of the US territory: *Acer rubrum* is a species of red maple that has made the fame of the canadian maple syrup, only located in the eastern part of North America ; *Toxicodendron diversilobum* or Pacific poison oak is a tree only present on the west coast and known for allergic reactions.

Geographic distribution of the tree categories:

Geographic distribution of tree categories across the US

Binned observation counts by foliage type.




We see different distribution of the tree categories across the US territory: - the most represented ones such as Deciduous Broadleaf have been observed all across the US territory - others, such as Deciduous Needleleaf have only been observed near the coasts and in southern parts of the US, namely because they easily lose their leaves, which makes the next spring more difficult for the tree in very cold areas.

Species richness by categories:

Species richness by tree category across the US

Number of distinct species per 1-degree grid cell, faceted by foliage type.



Once again, we see the large difference between the species of each category and how some of them are spread across the US.

11 Overall Interpretation

This report has examined the GlobalGeoTree dataset from multiple angles — source bias, temporal trends, taxonomic composition, worldwide distribution, regional patterns in Germany, and a country-level analysis of the United States.

Source bias and temporal trends. The global analysis confirmed that citizen-science platforms — above all iNaturalist — dominate the dataset and drive both its temporal growth and its geographic biases, making the combined dataset a composite of several geographically non-uniform “lenses” rather than a single coherent survey.

Taxonomic composition and worldwide distribution. The taxonomic analysis showed that Deciduous Broadleaf trees dominate observations globally, and that species richness peaks in temperate latitudes of both hemispheres. The continent-level analysis revealed that sampling effort is heavily concentrated in North America and Europe, while regions such as central Africa, the Amazon interior, and Central Asia remain sparsely sampled across the dataset years. Importantly, this geographic bias has remained structurally stable over 2015–2024, with growth taking the form of densification within already-sampled areas rather than expansion into new regions.

Germany and Berlin analysis. Within Germany, the same citizen-science dominance reproduces at the national level, with several platforms shaping partially different spatial footprints. At the Bundesland level, observation counts vary widely, reflecting differences in population density, platform coverage, and observer activity. Berlin is examined in detail, while Baden-Württemberg and NRW provide only brief map-based context for larger federal states with high observation counts. The Berlin oak analysis demonstrates a practical application of the dataset: identifying areas with high concentrations of recorded oak trees to support monitoring of oak processionary moth risk.

United States. The US analysis — the most-represented country in the dataset — showed that the most frequently recorded species (e.g. *Acer negundo*, *Acer rubrum*) correspond to genuinely common North American trees, but that species richness values are higher near the coasts where sampling effort is greatest. The central US shows lower recorded diversity, which may partly reflect lower sampling effort rather than genuinely lower biodiversity.

Limitations and outlook. All visualisations in this report show **recording intensity** in the dataset, not a direct census of trees. High values can reflect true tree presence, but they equally reflect where observers and platforms are most active. For any analysis built on GlobalGeoTree, the source composition of the sample should be reported alongside geographic, temporal, or taxonomic results, and key findings should be checked with and without the largest citizen-science source. Future work could refine the source classification rules, particularly for non-English source names, and explore whether the patterns observed in this analysis hold when the dominant citizen-science sources are excluded as a sensitivity check.

References

- Code for Germany. (n.d.). *Berlin district boundaries* [GeoJSON data file]. GitHub. https://github.com/codeforgermany/click_that_hood
- European Commission, Eurostat. (2021). *NUTS – Nomenclature of territorial units for statistics: NUTS 2021 boundary data* [GeoJSON data files]. GISCO — Geographical Information System of the Commission. <https://gisco-services.ec.europa.eu/distribution/v2/nuts/>
- GBIF Secretariat. (2024). *GBIF species API v1* [Application programming interface]. Global Biodiversity Information Facility. <https://www.gbif.org/developer/species>
- Mu, Y., Xiong, Z., Wang, Y., Shahzad, M., Essl, F., Kreft, H., van Kleunen, M., & Zhu, X. X. (2025). GlobalGeoTree: A multi-granular vision-language dataset for global tree species classification. *arXiv preprint arXiv:2505.12513*. <https://doi.org/10.48550/arXiv.2505.12513>